

### HPC unit 3

MPI (message passing interface) : a standardised and portable message passing system designed to function on a wide variety of parallel computing architectures. It enables the development of parallel programs that can run efficiently on both small clusters & large supercomputers.

It's a protocol used for programming parallel computers. Allows processes to communicate with one another by sending & receiving msgs.

MPI program structure:

MPI include file

Program begins



Initialise MPI environment (Parallel code)



Do work & make msg passing calls



Terminate MPI env.



Program ends.

MPI Communicators and Groups:

It uses objects called communicators & groups to define which collectn of processes may communicate with each other.

MPI-COMM-WORLD: default communicator

that includes all the processes initiated by the MPI program. Every process within MPI-comm-world can communicate with any other process in the same communicator.

level of thread support:

MPI-thread-single: level 0: only one thread will execute.

MPI-thread-funcctd (level 1): multiple threads may exist, but only main thread makes mpi calls.

MPI-thread-serialized (level 2): multiple threads can make mpi calls, but not concurrently; calls are serialized.

MPI-thread-multiple (level 3): multiple threads can make mpi calls without restrictions.

1) MPI-init: Initialize MPI executn environment  
Synopsis:

int MPI-init(int \*argc, char \*\*argv)

||P params:

argc: pointer to no. of args.

argv: pointer to the arg vector.

must be called by one thread only, that thread is called main thread & must be the thread that calls MPI-finalise.

2) MPI-Comm-Size: determines the size of grp associated with communicator.

int MPI-Comm-Size (in MPI-Comm comm, int \*size)

||P: communicator

o/p: no. of processes in the grp of comm int.

3) MPI-Comm-rank: determines the rank of calling processes in communicator

Synopsis: `int MPI-Comm-rank (MPI-Comm comm,  
int * rank)`

l/p: Comm

o/p: rank

4) MPI-Send: sends a msg to specified destination

l/p params:

- 1) buf: initial address of send buffer
- 2) count: no. of elements in " "
- 3) datatype: datatype of each send buffer elem
- 4) dest: rank of dest<sup>n</sup> process
- 5) tag: message tag
- 6) comm: comm handle.

5) MPI-Recv: receives msg from specified src.

l/p: same as mpi-send except buf.

o/p: buf: initial address of receive buffer

Status: status of received msg.

MPI-Scatter: distributes distinct chunks of data from the root process to all processes in a communicator.

Syntax:

`int MPI-Scatter (const void * sendbuf, int send count, MPI-Datatype sendtype, void * recvbuf, int recvcount, MPI-Datatype recvtype, int root, MPI-Comm comm)`

MPI-Gather: collects chunks of data from all processes and gathers them into a single buffer on root process.

Syntax same as above.

MPI-Bcast: broadcast data from the root process to all other processes in communicator

MPI-Barrier: blocks processes until all have reached this routine, synchronizing them